

# Introduction to HIP-Python



Norwegian research infrastructure services

**Hicham Agueny**  
**Scientific Computing Group**  
**University of Bergen/NRIS**

# Motivation

- **Python** is **widely used** in ML, data science, scientific computing
- Writing codes with **HIP C++** can be **challenging**
- Making GPU programming accessible to Python Users
- **Need of a python interface to interact with HIP libraries (C++ API)**

# Learning outcomes

- Access the GPU properties from a python interface
- Transfer data between host and device
- Call hipBLAS library from a python interface
- Compile and launch GPU Kernels from a python interface

**Examples:**

[https://github.com/HichamAgueny/HIP-Python\\_examples](https://github.com/HichamAgueny/HIP-Python_examples)

# Overview

## ❑ What is HIP-Python ?

## ❑ GPU management

- Get the device properties [hip.hipGetDeviceProperties](#)
- Get the device attributes [hip.hipDeviceGetAttribute](#)

## ❑ Memory management

- Allocate GPU memory [hip.hipMalloc\(\)](#)
- Free GPU memory [hip.hipFree\(\)](#)
- Direct memory copy [hip.hipMemcpy\(\)](#)

## ❑ Call HIP libraries from python

- Example: hipBLAS

## ❑ GPU Kernels

- HIP RTC API for compiling
- HIP runtime API functions for launching

# What is HIP-Python ?

HIP-Python is a **python wrapper** for:

- **HIP\* runtime** API (\*HIP-Heterogeneous-Compute Interface for Portability)
- **HIP RunTime Compilation (HIPRTC)** API
- Various **math libraries**, & communication library RCCL
- **Supports** both AMD and NVIDIA GPUs
  
- **HIP runtime API** provides functions to:
  - Access Device properties
  - Allocate and Free Device memory
  - Transfer Data between Host and Device
  - Launch Device Kernels (Device kernel is a function that is executed on the GPU)
  
- **HIPRTC API** provides functions to:
  - Create functions and managing GPU programs dynamically
  - Compile the Device Kernels at the runtime (i.e. during the execution of the application)

# Device properties via `hip.hipGetDeviceProperties()`

The object `hip.hipDeviceProp_t()` is passed as an argument to `hip.hipGetDeviceProperties()`

```
from hip import hip

props = hip.hipDeviceProp_t()
hip.hipGetDeviceProperties(props,0)
# Get selected properties

print(f"props.name = {props.name}")
print(f"props.gcnArchName = {props.gcnArchName}")
print(f"props.pciDeviceID = {props.pciDeviceID}")
print(f"props.totalGlobalMem = {props.totalGlobalMem}")
print(f"props.l2CacheSize = {props.l2CacheSize}")
```

**Out:**

```
props.name = b'AMD Instinct MI250X'
props.gcnArchName = b'gfx90a:sramecc+:xnack-'
props.pciDeviceID = 0
props.totalGlobalMem = 68702699520
props.l2CacheSize = 8388608
```

# Device properties via `hip.hipGetDeviceProperties()`

The object `hip.hipDeviceProp_t()` is passed as an argument to `hip.hipGetDeviceProperties()`

```
from hip import hip

props = hip.hipDeviceProp_t()
hip.hipGetDeviceProperties(props,0)

# Get all the device properties

for attrib in sorted(props.PROPERTIES()):
    print(f"props.{attrib}={getattr(props,attrib)}")
```

# Device properties via `hip.hipDeviceGetAttribute`

The outcome of `hip.hipDeviceAttribute_t()` is passed as an argument to `hip.hipDeviceGetAttribute()`

```
from hip import hip

for attrib in (
    hip.hipDeviceAttribute_t.hipDeviceAttributeMaxBlockDimX,
    hip.hipDeviceAttribute_t.hipDeviceAttributeMaxBlockDimY,
    hip.hipDeviceAttribute_t.hipDeviceAttributeMaxBlockDimZ,
    hip.hipDeviceAttribute_t.hipDeviceAttributeWarpSize,
    hip.hipDeviceAttribute_t.hipDeviceAttributeMaxThreadsPerBlock,
):
    result_attr = hip.hipDeviceGetAttribute(attrib, device_id)
    print(f"{attrib.name}: {result_attr[1]}")
```

**Out:**

- `hipDeviceAttributeMaxBlockDimX: 1024`
- `hipDeviceAttributeMaxBlockDimY: 1024`
- `hipDeviceAttributeMaxBlockDimZ: 1024`
- `hipDeviceAttributeWarpSize: 64`
- `hipDeviceAttributeMaxThreadsPerBlock: 1024`



# Memory management

- Allocate Device memory

**Ptr** = hip.hipMalloc(*unsigned long sizeBytes*)

- **Ptr**: Pointer to the memory to be allocated on the GPU
- **sizeBytes**: Data size in Bytes.

- Copy data from src (in) to **dst** (out)

**hip.hipMemcpy**(**dst**, **src**, *unsigned long sizeBytes*, **kind**)

- **dst**: This is the destination where the data will be copied to.
- **src**: This is the source from where the data will be copied.
- **kind**: Direction of transfer HostToDevice or DeviceToHost.

- Free Device memory

**hip.hipFree**(**Ptr**)

- **Ptr**: Pointer to the memory to be freed

# Memory management - Example 2

```
# Import some modules
import numpy as np
from hip import hip

# Generate random 1D-array
N = 10 #length
host_data = np.random.rand(N).astype(np.float32)

# Allocate device memory
num_bytes = N * np.float32().itemsize
device_data = hip.hipMalloc(num_bytes)

# Copy data from host to device
hip.hipMemcpy(device_data, host_data, num_bytes, hip.hipMemcpyKind.hipMemcpyHostToDevice)

# Copy data from device to host
host_data_b = np.empty_like(host_data)
hip.hipMemcpy(host_data_b, device_data, num_bytes, hip.hipMemcpyKind.hipMemcpyDeviceToHost)

# Free device memory
hip.hipFree(device_data)
```

# Calling HIP library from python

List of HIP libraries that can be called from a python interface through HIP-Python API.

- **hip.rccl**: Collective communication library (e.g. broadcast, reduce, ...) for multiple GPUs
- **hip.hiprtc**: HIP RunTime Compilation for compiling GPU-kernels (HIP C++) at runtime

## ➤ Math libraries

- **hip.hiprand**: Random number generation library optimized for AMD GPU
- **hip.hipfft**: Fast Fourier Transform library optimized for AMD GPU
- **hip.hipspase**: Sparse matrix operations library (sparse matrix-vector and matrix-matrix operations) optimized for AMD GPU
- **hip.hipsolver**: Dense linear algebra operations (solving linear systems) optimized for AMD GPU
- **hip.hipblas**: Basic Linear Algebra Subprograms (e.g. vector addition, matrix multiplication) optimized for AMD GPU

# Calling HipBLAS - Example 3

Calling a **hipBLAS function** from python interface consists of 3 steps

- Initiate hipBLAS
- Call a hipBLAS function to do computation on the GPU
- Destroy the library handle

**Example:** **hipblasSasum** is a function designed to compute the **sum of absolute values** of the elements in a single-precision.

- **Initiate hipBlas**

**Handle = hipblas.hipblasCreate()**

- **Call a hipblasSasum function**

**hipblas.hipblasSasum(handle, int n, x, int incx, result)**

- **Destroy handle**

**hipblas.hipblasDestroy(handle)**

- **handle:** This is the handle to the hipBLAS library context
- **n:** The number of elements in the input vector x
- **x:** The pointer to the input vector x storing the **n** elements
- **incx:** The increment between each element of x

- If incx is 1, the function will use each element of x.
- If incx is 2, the function will use every second element of x.

- **result:** The pointer to the variable where the result will be stored

# Take-away

- ❑ HIP-Python provides python users with a simple way to interact with HIP libraries and API runtimes.
- ❑ HIP libraries include:
  - **Math libraries:** hipBLAS, hipRAND, hipFFT, hipSPARSE, hipSOLVER.
  - **HIP Runtime API:** Allocate & free GPU memory, Copy data between Host and Device and launch GPU kernels.
  - **HIPRTC API:** Compile GPU kernels written with HIP C++.
- ❑ **Supports** both AMD and NVIDIA GPUs

# **Hands-on Examples**

# Hands-on examples

## 1-Download examples:

```
$ git clone https://github.com/HichamAgueny/HIP-Python\_examples.git
```

## 2-Launch an interactive session

```
$ salloc -A project_465001310 -t 00:30:00 -p dev-g -N 1 --gpus 1
```

## 3-Load modules

```
module load LUMI/24.03 partition/G
```

```
module load cray-python/3.11.7
```

## 4-Source the virtual env. where hip-python and numpy are installed

```
$ source /project/project_465001310/workshop_software/HIP-Python_examples/MyVirtEnv_hip_pytorch/bin/activate
```

## 5-Run examples

```
$ cd HIP-Python_examples
```


```
$ srun python example1_DeviceProp/deviceProp.py
```

```
$ srun python example2_MemoryManagemnt/memoryManagemnt.py
```

```
$ srun python example3_HipBlas/hipblasSum.py
```

```
$ srun python example4_DeviceKernels/deviceKernels.py
```

# GPU Kernel: Creating and compiling a program



## Compiling and launching GPU kernels:

### ➤ Creating and compiling a program with HIPRTC API

HIP allows to compile a program at runtime with HIPRTC API.

- Create the program `hiprtc.hiprtcCreateProgram()`
- Compile the program `hiprtc.hiprtcCompileProgram()`

### ➤ Launching a program with HIP runtime API

- Build module from object `hip.hipModuleLoadData()`
- Get the kernel function `hip.hipModuleGetFunction()`
- Specify grid and block dimensions `hip.dim3()`
- Launch the kernel `hip.hipModuleLaunchKernel()`



# GPU Kernel: Creating and compiling a program **Example 3**

- HIP allows to compile a program at runtime with HIPRTC API.
- Kernels can be stored as a text string and passed as an argument to `hiprtc.hiprtcCreateProgram()`
- The output from creating program is passed as input to `hiprtc.hiprtcCompileProgram()`

HIP C++ Kernel  
stored in a file  
named “kernel.hip”

```
extern "C" __global__ void Kernel_test(int n, float factor, float *A)
{
    int idx = threadIdx.x + blockIdx.x * blockDim.x;

    A[idx] = factor*A[idx];
}
```

Create program

```
prog =                                # hiprtc program hiprtc.hiprtcCreateProgram(kernel_source.encode()),#
Source code
                                b"Kernel_test",    # Name of Kernel
                                0,                # Number of Headers
                                [],                # Names of Headers
                                [])                # Names of Includes
```

Compile program

```
arch = b'gfx90a'

cflags = [b"--offload-arch="+arch]

hiprtc.hiprtcCompileProgram(prog, len(cflags), cflags)
```

# GPU Kernel: Launching a program **Example 4**

**Build module and  
Get the kernel**

```
code_size = hiprtc.hiprtcGetCodeSize(prog)
code = bytearray(code_size)
hiprtc.hiprtcGetCode(prog, code)
module = hip.hipModuleLoadData(code)
kernel = hip.hipModuleGetFunction(module, b"Kernel_test")
```

**Specify the block and  
grid dimensions**

```
block = hip.dim3(x=16, y=1, z=1)
grid = hip.dim3(math.ceil(N/block.x))
```

**Launch the program**

```
hip.hipModuleLaunchKernel(
    kernel,
    *grid,
    *block,
    sharedMemBytes=0,
    stream=None,
    kernelParams=None,
    extra=(
        ctypes.c_int(N),
        ctypes.c_float(factor),
        device_data,
    ))
```

# References

[https://rocm.docs.amd.com/projects/hip-python/en/latest/user\\_guide/1\\_usage.html#basic-usage-python](https://rocm.docs.amd.com/projects/hip-python/en/latest/user_guide/1_usage.html#basic-usage-python)

<https://github.com/ROCm/hip-python>

[https://github.com/HichamAgueny/HIP-Python\\_examples](https://github.com/HichamAgueny/HIP-Python_examples)